

Experiment - 3

1 **Aim:** Design a CNN architecture to implement the image classification task over an image dataset. Perform the Hyperparameter tuning and record the results.

2 Dataset Description

- The data that will be incorporated is the **MNIST database** which contains 60,000 images for training and 10,000 test images.
- The dataset consists of small square 28×28 pixel grayscale images of handwritten single digits between 0 and 9
- The MNIST dataset is conveniently bundled within Keras, and we can easily analyze some of its features in Python.

```
[1]: from tensorflow import keras
from keras.datasets import mnist      # MNIST dataset is included in Keras
(X_train, y_train), (X_test, y_test) = mnist.load_data()

print("X_train shape", X_train.shape)
print("y_train shape", y_train.shape)
print("X_test shape", X_test.shape)
print("y_test shape", y_test.shape)
```

```
WARNING:tensorflow:From c:\MAGNUS\files\GAI\aiml\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```


```
X_train shape (60000, 28, 28)
y_train shape (60000,)
X_test shape (10000, 28, 28)
y_test shape (10000,)
```

```
[2]: # Visualize any random image
# Plot first few images
import matplotlib.pyplot as plt
```

```

for i in range(9):
    # define subplot
    plt.subplot(3,3,i+1) # 3 rows, 3 col, pos
    # plot raw pixel data
    plt.imshow(X_train[i],cmap='gray')
# show the figure
plt.show()

```



2.0.1 Formatting the Input

```

[3]: # Single-channel input data (grey-scale)
      # First apply convolutions then flatten

X_train = X_train.reshape(60000, 28, 28, 1) # single-channel input
X_test = X_test.reshape(10000, 28, 28, 1)

X_train = X_train.astype('float32')           # change integers to 32-bit
      ↪ floating point numbers
X_test = X_test.astype('float32')

X_train /= 255                                # min-max normalization

```

```
X_test /= 255

print("Training matrix shape", X_train.shape)
print("Testing matrix shape", X_test.shape)
```

Training matrix shape (60000, 28, 28, 1)

Testing matrix shape (10000, 28, 28, 1)

3 Convolutional Neural Network

- Convolution applies **kernels** (filters) that traverse through each image and generate **feature maps**
- keras Conv2D: https://keras.io/api/layers/convolution_layers/convolution2d/
- Each kernel in a CNN learns a different characteristic of an image.
- **max pooling** helps in reducing the number of learnable parameters, and decreasing the computational cost (e.g. system memory)

3.1 Building a Convolutional Neural Network

```
[4]: from keras import backend as K
from keras import __version__

print('Using Keras version:', __version__, 'backend:', K.backend())
```

Using Keras version: 2.15.0 backend: tensorflow

```
[5]: # import cnn layers
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, Dense
import tensorflow as tf
```

```
[6]: model = Sequential()                                # Linear stacking of layers

# Convolution Layer 1: 8 filters, kernel size 3x3, relu activation, valid
# padding, stride 1
model.add(Conv2D(8, kernel_size=(3,3), activation='relu', strides=1,
# padding='valid', input_shape=(28, 28, 1)))
# MaxPooling: pool size 2, stride 2
model.add(MaxPooling2D(pool_size=2, strides=2))
# Convolution Layer 2: 16 filters, kernel size 3x3, relu activation, valid
# padding, stride 1
model.add(Conv2D(16, kernel_size=(3,3), activation='relu', strides=1,
# padding='valid'))
# MaxPooling: pool size 2, stride 2
model.add(MaxPooling2D(pool_size=2, strides=2))
# Flatten final feature matrix into a 1d array
model.add(Flatten(input_shape=(28,28)))
```

```

# Fully Connected Layer: 64 units and relu activation
model.add(Dense(64,activation='relu'))
# Dropout layer, 0.2 rate
model.add(Dropout(rate=0.2))
# Final output dense Layer
model.add(Dense(10, activation='softmax'))

#Compile the model with sparse_categorical_crossentropy loss
model.
↪compile(loss='sparse_categorical_crossentropy',optimizer='adam',metrics=['accuracy'])

```

WARNING:tensorflow:From c:\MAGNUS\files\GAI\aiml\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From c:\MAGNUS\files\GAI\aiml\Lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

WARNING:tensorflow:From c:\MAGNUS\files\GAI\aiml\Lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

```

[7]: model.build()
      model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d (MaxPooling2D)	(None, 13, 13, 8)	0
conv2d_1 (Conv2D)	(None, 11, 11, 16)	1168
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten (Flatten)	(None, 400)	0
dense (Dense)	(None, 64)	25664
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 10)	650

```
=====
Total params: 27562 (107.66 KB)
Trainable params: 27562 (107.66 KB)
Non-trainable params: 0 (0.00 Byte)
-----
```

```
[8]: # Conv1: 3x3 kernels, one for each the single channel, 8 such filters and 8
      ↪biases
print('Conv1: ',3*3*1*8 + 8)
# Conv2: 3x3 kernels, one for each of the 8 channels, 16 such filters and 16
      ↪biases
print('Conv2: ',3*3*8*16 + 16)
# input to dense layer
print('Flatten:', 5*5*16)
# 400 inputs, 1 bias connected to each of 64 units in dense layer
print('Dense1: ',400*64+64)
# 64 inputs, 1 bias connected to each of 10 units in output layer
print('Dense2: ',64*10+10)
```

```
Conv1: 80
Conv2: 1168
Flatten: 400
Dense1: 25664
Dense2: 650
```

```
[9]: # Visualize the model
from keras.utils import plot_model
plot_model(model, show_shapes=True, show_layer_names=False)
```

```
[9]:
```

InputLayer	input:	[(None, 28, 28, 1)]
	output:	[(None, 28, 28, 1)]



Conv2D	input:	(None, 28, 28, 1)
	output:	(None, 26, 26, 8)



MaxPooling2D	input:	(None, 26, 26, 8)
	output:	(None, 13, 13, 8)



Conv2D	input:	(None, 13, 13, 8)
	output:	(None, 11, 11, 16)



MaxPooling2D	input:	(None, 11, 11, 16)
	output:	(None, 5, 5, 16)



Flatten	input:	(None, 5, 5, 16)
	output:	(None, 400)



Dense	input:	(None, 400)
	output:	(None, 64)



Dropout	input:	(None, 64)
	output:	(None, 64)



Dense	input:	(None, 64)
	output:	(None, 10)

Train the model

- Validation data = $0.2 \times 60,000 = 12,000$
- Batch size = 128
- Number of batches during training are $(60000-12000)/128 = 48000/128 = 375$

```
[10]: # Train the model
batch_size=128
epochs=10
hist = model.fit(X_train,
    ↪ y_train, epochs=epochs, batch_size=batch_size, verbose=1, validation_split=0.2)
```

Epoch 1/10

WARNING:tensorflow:From c:\MAGNUS\files\GAI\aiml\Lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From c:\MAGNUS\files\GAI\aiml\Lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

375/375 [=====] - 5s 7ms/step - loss: 0.5493 - accuracy: 0.8431 - val_loss: 0.1316 - val_accuracy: 0.9622

Epoch 2/10

375/375 [=====] - 2s 6ms/step - loss: 0.1364 - accuracy: 0.9595 - val_loss: 0.0811 - val_accuracy: 0.9765

Epoch 3/10

375/375 [=====] - 2s 6ms/step - loss: 0.0980 - accuracy: 0.9702 - val_loss: 0.0675 - val_accuracy: 0.9797

Epoch 4/10

375/375 [=====] - 2s 6ms/step - loss: 0.0806 - accuracy: 0.9753 - val_loss: 0.0629 - val_accuracy: 0.9816

Epoch 5/10

375/375 [=====] - 3s 8ms/step - loss: 0.0676 - accuracy: 0.9788 - val_loss: 0.0538 - val_accuracy: 0.9852

Epoch 6/10

375/375 [=====] - 3s 8ms/step - loss: 0.0591 - accuracy: 0.9822 - val_loss: 0.0495 - val_accuracy: 0.9861

Epoch 7/10

375/375 [=====] - 3s 8ms/step - loss: 0.0535 - accuracy: 0.9833 - val_loss: 0.0509 - val_accuracy: 0.9857

Epoch 8/10

375/375 [=====] - 3s 8ms/step - loss: 0.0470 - accuracy: 0.9856 - val_loss: 0.0479 - val_accuracy: 0.9870

Epoch 9/10

```
375/375 [=====] - 3s 7ms/step - loss: 0.0424 -  
accuracy: 0.9867 - val_loss: 0.0448 - val_accuracy: 0.9877  
Epoch 10/10  
375/375 [=====] - 2s 6ms/step - loss: 0.0403 -  
accuracy: 0.9876 - val_loss: 0.0434 - val_accuracy: 0.9884
```

3.1.1 Evaluate Model

```
[11]: score = model.evaluate(X_test, y_test, verbose = 0)  
print('Test loss:', score[0])  
print('Test accuracy:', score[1])
```

```
Test loss: 0.03615253418684006  
Test accuracy: 0.9889000058174133
```

```
[12]: # make one prediction  
print('Actual class:', y_test[0])  
print('Class Probabilities:')  
model.predict(X_test[0].reshape(1, 28, 28, 1))
```

```
Actual class: 7  
Class Probabilities:  
1/1 [=====] - 0s 103ms/step
```

```
[12]: array([[2.8043439e-06, 1.2572475e-07, 4.6315331e-06, 1.4247337e-05,  
          9.9764907e-10, 6.4714554e-09, 8.9647369e-12, 9.9997783e-01,  
          1.2663826e-08, 3.6554039e-07]], dtype=float32)
```

```
[13]: import numpy as np  
yhat_test = np.argmax(model.predict(X_test), axis=-1)  
print(yhat_test[0:10])  
print(y_test[0:10])
```

```
313/313 [=====] - 1s 2ms/step  
[7 2 1 0 4 1 4 9 5 9]  
[7 2 1 0 4 1 4 9 5 9]
```

```
[14]: from sklearn.metrics import accuracy_score  
print('Accuracy:')  
print(float(accuracy_score(y_test, yhat_test))*100, '%')
```

```
Accuracy:  
98.89 %
```

```
[15]: from sklearn.metrics import confusion_matrix  
print('Confusion Matrix:')  
print(confusion_matrix(y_test, yhat_test))
```

```
Confusion Matrix:  
[[ 977    0    0    0    1    0    0    1    1    0]
```




```
[ 0 1132 1 0 0 0 1 0 1 0]
[ 2 2 1022 0 2 0 0 2 2 0]
[ 1 0 0 997 0 4 0 1 5 2]
[ 0 0 0 0 978 0 0 0 1 3]
[ 2 0 1 5 0 879 2 1 1 1]
[ 5 2 0 0 3 2 942 0 4 0]
[ 0 3 8 1 1 0 0 1009 0 6]
[ 3 0 0 1 1 1 0 2 962 4]
[ 3 3 0 0 8 1 0 0 3 991]]
```


3.1.2 Plot Learning curves

```
[16]: hist.history.keys()
epochRange = range(1,epochs+1)
```

```
[17]: # Plot Accuracy vs epochs (DIY)
plt.plot(epochRange,hist.history['accuracy'])
plt.plot(epochRange,hist.history['val_accuracy'])
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.grid()
plt.xlim((1,epochs))
plt.legend(['Train','Test'])
plt.show()
```



```
[18]: # Plot Loss vs epochs (DIY)
plt.plot(epochRange,hist.history['loss'])
plt.plot(epochRange,hist.history['val_loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid()
plt.xlim((1,epochs))
plt.legend(['Train', 'Test'])
plt.show()
```



3.2 Removing the dropout layer

```
[19]: model_2 = Sequential() # Linear stacking of layers

# Convolution Layer 1: 8 filters, kernel size 3x3, relu activation, valid
padding, stride 1
model_2.add(Conv2D(8,kernel_size=(3,3),activation='relu',strides=1,
padding='valid',input_shape=(28, 28, 1)))
```

```

# MaxPooling: pool size 2, stride 2
model_2.add(MaxPooling2D(pool_size=2, strides=2))
# Convolution Layer 2: 16 filters, kernel size 3x3, relu activation, valid
↳padding, stride 1
model_2.add(Conv2D(16, kernel_size=(3,3), activation='relu', strides=1,
↳padding='valid'))
# MaxPooling: pool size 2, stride 2
model_2.add(MaxPooling2D(pool_size=2, strides=2))
# Flatten final feature matrix into a 1d array
model_2.add(Flatten(input_shape=(28,28)))
# Fully Connected Layer: 64 units and relu activation
model_2.add(Dense(64, activation='relu'))
# Final output dense Layer
model_2.add(Dense(10, activation='softmax'))
# Compile the new model with sparse_categorical_crossentropy loss
model_2.
↳compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

```

[20]: model_2.build()
      model_2.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 8)	0
conv2d_3 (Conv2D)	(None, 11, 11, 16)	1168
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten_1 (Flatten)	(None, 400)	0
dense_2 (Dense)	(None, 64)	25664
dense_3 (Dense)	(None, 10)	650
Total params: 27562 (107.66 KB)		
Trainable params: 27562 (107.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
[21]: plot_model(model_2, show_shapes=True, show_layer_names=False)
```

```
[21]:
```

InputLayer	input:	[(None, 28, 28, 1)]
	output:	[(None, 28, 28, 1)]



Conv2D	input:	(None, 28, 28, 1)
	output:	(None, 26, 26, 8)



MaxPooling2D	input:	(None, 26, 26, 8)
	output:	(None, 13, 13, 8)



Conv2D	input:	(None, 13, 13, 8)
	output:	(None, 11, 11, 16)



MaxPooling2D	input:	(None, 11, 11, 16)
	output:	(None, 5, 5, 16)



Flatten	input:	(None, 5, 5, 16)
	output:	(None, 400)



Dense	input:	(None, 400)
	output:	(None, 64)



Dense	input:	(None, 64)
	output:	(None, 10)

```
[22]: batch_size=128
epochs=10
hist_2 = model_2.fit(X_train,
    ↪ y_train, epochs=epochs, batch_size=batch_size, verbose=1, validation_split=0.2)
```

```
Epoch 1/10
375/375 [=====] - 5s 8ms/step - loss: 0.4212 -
accuracy: 0.8790 - val_loss: 0.1563 - val_accuracy: 0.9528
Epoch 2/10
375/375 [=====] - 3s 7ms/step - loss: 0.1156 -
accuracy: 0.9651 - val_loss: 0.1001 - val_accuracy: 0.9708
Epoch 3/10
375/375 [=====] - 3s 7ms/step - loss: 0.0839 -
accuracy: 0.9741 - val_loss: 0.0988 - val_accuracy: 0.9704
Epoch 4/10
375/375 [=====] - 2s 6ms/step - loss: 0.0673 -
accuracy: 0.9788 - val_loss: 0.0672 - val_accuracy: 0.9802
Epoch 5/10
375/375 [=====] - 2s 6ms/step - loss: 0.0564 -
accuracy: 0.9821 - val_loss: 0.0659 - val_accuracy: 0.9797
Epoch 6/10
375/375 [=====] - 2s 6ms/step - loss: 0.0472 -
accuracy: 0.9851 - val_loss: 0.0608 - val_accuracy: 0.9828
Epoch 7/10
375/375 [=====] - 2s 6ms/step - loss: 0.0443 -
accuracy: 0.9861 - val_loss: 0.0540 - val_accuracy: 0.9838
Epoch 8/10
375/375 [=====] - 2s 6ms/step - loss: 0.0375 -
accuracy: 0.9878 - val_loss: 0.0580 - val_accuracy: 0.9837
Epoch 9/10
375/375 [=====] - 2s 6ms/step - loss: 0.0343 -
accuracy: 0.9894 - val_loss: 0.0619 - val_accuracy: 0.9816
Epoch 10/10
375/375 [=====] - 2s 6ms/step - loss: 0.0303 -
accuracy: 0.9905 - val_loss: 0.0502 - val_accuracy: 0.9853
```

```
[37]: score_2 = model_2.evaluate(X_test, y_test, verbose = 0)
print('Test loss:', score_2[0])
print('Test accuracy:', score_2[1])
```

```
Test loss: 0.034381765872240067
Test accuracy: 0.9879999756813049
```

```
[48]: yhat_test = np.argmax(model_2.predict(X_test), axis=-1)
print(yhat_test[0:10])
```

```

print(y_test[0:10])
print('Confusion Matrix:')
print(confusion_matrix(y_test, yhat_test))
print('Accuracy:')
print(float(accuracy_score(y_test, yhat_test))*100,'%')

```

313/313 [=====] - 1s 3ms/step

[7 2 1 0 4 1 4 9 5 9]

[7 2 1 0 4 1 4 9 5 9]

Confusion Matrix:

```

[[ 972    0    0    0    0    3    2    3    0]
 [   0 1131    1    0    0    1    0    2    0]
 [   1    3 1021    0    1    1    1    3    0]
 [   1    0    1 1000    0    3    0    1    4    0]
 [   1    1    2    0 966    0    1    1    0 10]
 [   1    0    0    4    0 881    2    1    3    0]
 [   3    3    0    0    4    3 944    0    1    0]
 [   0    1   11    3    0    0    0 1009    1    3]
 [   2    0    3    1    1    0    0    1 964    2]
 [   1    2    0    2    5    2    0    5    0 992]]

```


Accuracy:

98.8 %


```

[25]: plt.plot(epochRange,hist_2.history['val_accuracy'])
plt.plot(epochRange,hist_2.history['accuracy'])
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.grid()
plt.xlim((1,epochs))
plt.legend(['Train','Test'])
plt.show()

```



```
[26]: plt.plot(epochRange,hist_2.history['loss'])
plt.plot(epochRange,hist_2.history['val_loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid()
plt.xlim((1,epochs))
plt.legend(['Train','Test'])
plt.show()
```

3.3 Increasing filters in Convolution layers

```
[38]: model_3 = Sequential() # Linear stacking of layers

# Convolution Layer 1: 64 filters, kernel size 3x3, relu activation, valid
padding, stride 1
model_3.add(Conv2D(64,kernel_size=(3,3),activation='relu',strides=1,
padding='valid',input_shape=(28, 28, 1)))
# MaxPooling: pool size 2, stride 2
model_3.add(MaxPooling2D(pool_size=2,strides=2))
# Convolution Layer 2: 64 filters, kernel size 3x3, relu activation, valid
padding, stride 1
model_3.add(Conv2D(64,kernel_size=(3,3),activation='relu',strides=1,
padding='valid'))
# MaxPooling: pool size 2, stride 2
model_3.add(MaxPooling2D(pool_size=2,strides=2))
# Flatten final feature matrix into a 1d array
model_3.add(Flatten(input_shape=(28,28)))
# Fully Connected Layer: 64 units and relu activation
model_3.add(Dense(64,activation='relu'))
```

```

# Dropout layer, 0.2 rate
model_3.add(Dropout(rate=0.2))
# Final output dense Layer
model_3.add(Dense(10, activation='softmax'))

#Compile the model with sparse_categorical_crossentropy loss
model_3.
    ↪compile(loss='sparse_categorical_crossentropy',optimizer='adam',metrics=['accuracy'])

```

```

[39]: model_3.build()
      model_3.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_6 (MaxPooling2D)	(None, 13, 13, 64)	0
conv2d_7 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_7 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_3 (Flatten)	(None, 1600)	0
dense_6 (Dense)	(None, 64)	102464
dropout_2 (Dropout)	(None, 64)	0
dense_7 (Dense)	(None, 10)	650
Total params: 140682 (549.54 KB)		
Trainable params: 140682 (549.54 KB)		
Non-trainable params: 0 (0.00 Byte)		

```

[40]: plot_model(model_3, show_shapes=True, show_layer_names=False)

```

[40]:

InputLayer	input:	[(None, 28, 28, 1)]
	output:	[(None, 28, 28, 1)]



Conv2D	input:	(None, 28, 28, 1)
	output:	(None, 26, 26, 64)



MaxPooling2D	input:	(None, 26, 26, 64)
	output:	(None, 13, 13, 64)



Conv2D	input:	(None, 13, 13, 64)
	output:	(None, 11, 11, 64)



MaxPooling2D	input:	(None, 11, 11, 64)
	output:	(None, 5, 5, 64)



Flatten	input:	(None, 5, 5, 64)
	output:	(None, 1600)



Dense	input:	(None, 1600)
	output:	(None, 64)



Dropout	input:	(None, 64)
	output:	(None, 64)



Dense	input:	(None, 64)
	output:	(None, 10)

```
[41]: batch_size=128
epochs=10
hist_3 = model_3.fit(X_train,
    ↪ y_train, epochs=epochs, batch_size=batch_size, verbose=1, validation_split=0.2)
```

Epoch 1/10

375/375 [=====] - 13s 31ms/step - loss: 0.2943 - accuracy: 0.9091 - val_loss: 0.0794 - val_accuracy: 0.9766

Epoch 2/10

375/375 [=====] - 11s 29ms/step - loss: 0.0883 - accuracy: 0.9733 - val_loss: 0.0597 - val_accuracy: 0.9829

Epoch 3/10

375/375 [=====] - 11s 28ms/step - loss: 0.0636 - accuracy: 0.9806 - val_loss: 0.0492 - val_accuracy: 0.9852

Epoch 4/10

375/375 [=====] - 11s 28ms/step - loss: 0.0505 - accuracy: 0.9842 - val_loss: 0.0441 - val_accuracy: 0.9870

Epoch 5/10

375/375 [=====] - 11s 28ms/step - loss: 0.0423 - accuracy: 0.9864 - val_loss: 0.0385 - val_accuracy: 0.9888

Epoch 6/10

375/375 [=====] - 11s 28ms/step - loss: 0.0350 - accuracy: 0.9891 - val_loss: 0.0368 - val_accuracy: 0.9892

Epoch 7/10

375/375 [=====] - 11s 29ms/step - loss: 0.0287 - accuracy: 0.9908 - val_loss: 0.0444 - val_accuracy: 0.9883

Epoch 8/10

375/375 [=====] - 11s 28ms/step - loss: 0.0266 - accuracy: 0.9913 - val_loss: 0.0383 - val_accuracy: 0.9894

Epoch 9/10

375/375 [=====] - 10s 27ms/step - loss: 0.0215 - accuracy: 0.9931 - val_loss: 0.0373 - val_accuracy: 0.9896

Epoch 10/10

375/375 [=====] - 10s 28ms/step - loss: 0.0205 - accuracy: 0.9931 - val_loss: 0.0422 - val_accuracy: 0.9900

```
[42]: score_3 = model_3.evaluate(X_test, y_test, verbose = 0)
print('Test loss:', score_3[0])
print('Test accuracy:', score_3[1])
```

Test loss: 0.03154883161187172

Test accuracy: 0.9904999732971191

```
[44]: yhat_test = np.argmax(model_3.predict(X_test), axis=-1)
print(yhat_test[0:10])
```

```

print(y_test[0:10])
print('Confusion Matrix:')
print(confusion_matrix(y_test, yhat_test))

```

313/313 [=====] - 1s 4ms/step

[7 2 1 0 4 1 4 9 5 9]

[7 2 1 0 4 1 4 9 5 9]

Confusion Matrix:

```

[[ 977    0    1    0    0    0    0    1    1    0]
 [   0 1131    0    2    0    0    0    0    2    0]
 [   2    0 1021    3    0    0    0    5    1    0]
 [   0    0    1 1006    0    1    0    0    2    0]
 [   1    1    0    0 968    0    5    1    1    5]
 [   2    0    0   10    0 871    1    1    6    1]
 [   5    3    0    1    1    1 941    0    6    0]
 [   0    1    3    0    0    0    0 1022    0    2]
 [   1    1    0    2    0    0    0    2 968    0]
 [   2    2    0    0    1    1    1    1    1 1000]]

```

```

[47]: print('Accuracy:')
      print(float(accuracy_score(y_test, yhat_test))*100,'%')

```


Accuracy:

99.05000000000001 %


```

[45]: plt.plot(epochRange,hist_3.history['accuracy'])
      plt.plot(epochRange,hist_3.history['val_accuracy'])
      plt.xlabel('Epoch')
      plt.ylabel('Accuracy')
      plt.grid()
      plt.xlim((1,epochs))
      plt.legend(['Train','Test'])
      plt.show()

```



```
[46]: plt.plot(epochRange,hist_3.history['loss'])
plt.plot(epochRange,hist_3.history['val_loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid()
plt.xlim((1,epochs))
plt.legend(['Train','Test'])
plt.show()
```



3.4 Changing kernel-size, stride and padding in Convolutional layer

```
[49]: model_4 = Sequential() # Linear stacking of layers

# Convolution Layer 1: 8 filters, kernel size 2x2, relu activation, same
padding, stride 2
model_4.add(Conv2D(8,kernel_size=(2,2),activation='relu',strides=2,
padding='same',input_shape=(28, 28, 1)))
# MaxPooling: pool size 2, stride 2
model_4.add(MaxPooling2D(pool_size=2,strides=2))
# Convolution Layer 2: 16 filters, kernel size 2x2, relu activation, same
padding, stride 2
model_4.add(Conv2D(16,kernel_size=(2,2),activation='relu',strides=2,
padding='same'))
# MaxPooling: pool size 2, stride 2
model_4.add(MaxPooling2D(pool_size=2,strides=2))
# Flatten final feature matrix into a 1d array
model_4.add(Flatten(input_shape=(28,28)))
# Fully Connected Layer: 64 units and relu activation
model_4.add(Dense(64,activation='relu'))
```

```

# Dropout layer, 0.2 rate
model_4.add(Dropout(rate=0.2))
# Final output dense Layer
model_4.add(Dense(10, activation='softmax'))

#Compile the model with sparse_categorical_crossentropy loss
model_4.
    ↪compile(loss='sparse_categorical_crossentropy',optimizer='adam',metrics=['accuracy'])

```

```

[50]: model_4.build()
      model_4.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 14, 14, 8)	40
max_pooling2d_8 (MaxPooling2D)	(None, 7, 7, 8)	0
conv2d_9 (Conv2D)	(None, 4, 4, 16)	528
max_pooling2d_9 (MaxPooling2D)	(None, 2, 2, 16)	0
flatten_4 (Flatten)	(None, 64)	0
dense_8 (Dense)	(None, 64)	4160
dropout_3 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 10)	650
Total params: 5378 (21.01 KB)		
Trainable params: 5378 (21.01 KB)		
Non-trainable params: 0 (0.00 Byte)		

```

[51]: plot_model(model_4, show_shapes=True, show_layer_names=False)

```

[51]:

InputLayer	input:	[(None, 28, 28, 1)]
	output:	[(None, 28, 28, 1)]



Conv2D	input:	(None, 28, 28, 1)
	output:	(None, 14, 14, 8)



MaxPooling2D	input:	(None, 14, 14, 8)
	output:	(None, 7, 7, 8)



Conv2D	input:	(None, 7, 7, 8)
	output:	(None, 4, 4, 16)



MaxPooling2D	input:	(None, 4, 4, 16)
	output:	(None, 2, 2, 16)



Flatten	input:	(None, 2, 2, 16)
	output:	(None, 64)



Dense	input:	(None, 64)
	output:	(None, 64)



Dropout	input:	(None, 64)
	output:	(None, 64)



Dense	input:	(None, 64)
	output:	(None, 10)

```
[52]: batch_size=128
epochs=10
hist_4 = model_4.fit(X_train,
    ↪ y_train, epochs=epochs, batch_size=batch_size, verbose=1, validation_split=0.2)
```

Epoch 1/10

375/375 [=====] - 4s 6ms/step - loss: 1.1088 - accuracy: 0.6406 - val_loss: 0.5164 - val_accuracy: 0.8395

Epoch 2/10

375/375 [=====] - 2s 5ms/step - loss: 0.5424 - accuracy: 0.8238 - val_loss: 0.3939 - val_accuracy: 0.8745

Epoch 3/10

375/375 [=====] - 2s 6ms/step - loss: 0.4458 - accuracy: 0.8585 - val_loss: 0.3352 - val_accuracy: 0.8943

Epoch 4/10

375/375 [=====] - 2s 6ms/step - loss: 0.3961 - accuracy: 0.8744 - val_loss: 0.2988 - val_accuracy: 0.9057

Epoch 5/10

375/375 [=====] - 2s 6ms/step - loss: 0.3610 - accuracy: 0.8871 - val_loss: 0.2859 - val_accuracy: 0.9097

Epoch 6/10

375/375 [=====] - 2s 5ms/step - loss: 0.3336 - accuracy: 0.8934 - val_loss: 0.2621 - val_accuracy: 0.9184

Epoch 7/10

375/375 [=====] - 2s 5ms/step - loss: 0.3118 - accuracy: 0.9010 - val_loss: 0.2439 - val_accuracy: 0.9230

Epoch 8/10

375/375 [=====] - 2s 4ms/step - loss: 0.2958 - accuracy: 0.9053 - val_loss: 0.2286 - val_accuracy: 0.9288

Epoch 9/10

375/375 [=====] - 2s 4ms/step - loss: 0.2810 - accuracy: 0.9103 - val_loss: 0.2243 - val_accuracy: 0.9302

Epoch 10/10

375/375 [=====] - 2s 5ms/step - loss: 0.2652 - accuracy: 0.9161 - val_loss: 0.2110 - val_accuracy: 0.9352

```
[53]: score_4 = model_4.evaluate(X_test, y_test, verbose = 0)
print('Test loss:', score_4[0])
print('Test accuracy:', score_4[1])
```

Test loss: 0.20662404596805573

Test accuracy: 0.9318000078201294

```
[54]: yhat_test = np.argmax(model_4.predict(X_test), axis=-1)
print(yhat_test[0:10])
```

```

print(y_test[0:10])
print('Confusion Matrix:')
print(confusion_matrix(y_test, yhat_test))

```

313/313 [=====] - 1s 3ms/step

[7 2 1 0 4 1 4 9 5 9]

[7 2 1 0 4 1 4 9 5 9]

Confusion Matrix:

```

[[ 947    1    4    3    0    5   11    5    4    0]
 [   0 1110    7    5    0    1    5    2    5    0]
 [  14    1  973   10    2    0    5   17   10    0]
 [   2    2   21  916    0   36    2    9   20    2]
 [   3    2    2    1  911    2   12    6    4   39]
 [   8    0    1   39    0  816   17    2    4    5]
 [  15    2    5    0    3   13  919    0    1    0]
 [   3    5   19    6    8    0    0  964    6   17]
 [  10    1    8   21    2   27   10    4  879   12]
 [   6    7    4    8   58   17    3   17    6  883]]

```

```

[55]: print('Accuracy:')
      print(float(accuracy_score(y_test, yhat_test))*100,'%')

```


Accuracy:

93.17999999999999 %


```

[56]: plt.plot(epochRange,hist_4.history['accuracy'])
      plt.plot(epochRange,hist_4.history['val_accuracy'])
      plt.xlabel('Epoch')
      plt.ylabel('Accuracy')
      plt.grid()
      plt.xlim((1,epochs))
      plt.legend(['Train','Test'])
      plt.show()

```



```
[57]: plt.plot(epochRange,hist_4.history['loss'])
plt.plot(epochRange,hist_4.history['val_loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid()
plt.xlim((1,epochs))
plt.legend(['Train','Test'])
plt.show()
```



4 Conclusion

A Convolutional Neural Network architecture was designed to implement the image classification task over the MNIST image dataset. Hyper-parameter tuning was performed and results were recorded.

The model with dropout rate set to 0.2 has the accuracy 98.89%.

The model with no dropout layer has the accuracy 98.79%.

The model with increased filters in Conv2D layers has the accuracy 99.05%.

The model with kernel-size set to (2,2), stride = 2 and same padding has the accuracy 93.18%.